

# Practical\_Machine\_Learning

Victor Alanis

28/7/2020

## Practical Machine Learning Project : Prediction Assignment Writeup

### I. Overview

This document is the final report of the Peer Assessment project from Coursera's course Practical Machine Learning, as part of the Specialization in Data Science. It was built up in RStudio, using its knitr functions, meant to be published in html format. This analysis meant to be the basis for the course quiz and a prediction assignment writeup. The main goal of the project is to predict the manner in which 6 participants performed some exercise as described below. This is the "classe" variable in the training set. The machine learning algorithm described here is applied to the 20 test cases available in the test data and the predictions are submitted in appropriate format to the Course Project Prediction Quiz for automated grading.

### II. Background

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement - a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. More information is available from the website here: <http://groupware.les.inf.puc-rio.br/har>(see the section on the Weight Lifting Exercise Dataset).

Read more:<http://groupware.les.inf.puc-rio.br/har#ixzz3xsbs5bVX>

### III. Data Loading and Exploratory Analysis

**Environment Preparation** We first upload the R libraries that are necessary for the complete analysis.

```
## Warning: package 'knitr' was built under R version 4.0.2
```

```
## Loading required package: lattice
```

```
## Warning: package 'lattice' was built under R version 4.0.2
```

```
## Loading required package: ggplot2
```

```
## Warning: package 'ggplot2' was built under R version 4.0.2
```

```
## Warning: package 'rpart.plot' was built under R version 4.0.2

## Warning: package 'rattle' was built under R version 4.0.2

## Loading required package: tibble

## Loading required package: bitops

## Rattle: A free graphical interface for data science with R.
## Versión 5.4.0 Copyright (c) 2006-2020 Togaware Pty Ltd.
## Escriba 'rattle()' para agitar, sacudir y rotar sus datos.

## Warning: package 'randomForest' was built under R version 4.0.2

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:rattle':
##
##     importance

## The following object is masked from 'package:ggplot2':
##
##     margin

## Warning: package 'corrplot' was built under R version 4.0.2

## corrplot 0.84 loaded
```

**Data Loading and Cleaning** The next step is loading the dataset from the URL provided above. The training dataset is then partitioned in 2 to create a Training set (70% of the data) for the modeling process and a Test set (with the remaining 30%) for the validations. The testing dataset is not changed and will only be used for the quiz results generation.

```
# set the URL for the download
UrlTrain <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"
UrlTest  <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"

# download the datasets
training <- read.csv(url(UrlTrain))
testing  <- read.csv(url(UrlTest))

# create a partition with the training dataset
inTrain  <- createDataPartition(training$classe, p=0.7, list=FALSE)
TrainSet <- training[inTrain, ]
TestSet  <- training[-inTrain, ]
dim(TrainSet)

## [1] 13737 160
```

```
dim(TestSet)
```

```
## [1] 5885 160
```

Both created datasets have 160 variables. Those variables have plenty of NA, that can be removed with the cleaning procedures below. The Near Zero variance (NZV) variables are also removed and the ID variables as well.

```
# remove variables with Nearly Zero Variance
```

```
NZV <- nearZeroVar(TrainSet)
```

```
TrainSet <- TrainSet[, -NZV]
```

```
TestSet <- TestSet[, -NZV]
```

```
dim(TrainSet)
```

```
## [1] 13737 104
```

```
dim(TestSet)
```

```
## [1] 5885 104
```

```
# remove variables that are mostly NA
```

```
AllNA <- sapply(TrainSet, function(x) mean(is.na(x))) > 0.95
```

```
TrainSet <- TrainSet[, AllNA==FALSE]
```

```
TestSet <- TestSet[, AllNA==FALSE]
```

```
dim(TrainSet)
```

```
## [1] 13737 59
```

```
dim(TestSet)
```

```
## [1] 5885 59
```

```
# remove identification only variables (columns 1 to 5)
```

```
TrainSet <- TrainSet[, -(1:5)]
```

```
TestSet <- TestSet[, -(1:5)]
```


```
dim(TrainSet)
```

```
## [1] 13737 54
```

**Correlation Analysis** A correlation among variables is analysed before proceeding to the modeling procedures.

```
corMatrix <- cor(TrainSet[, -54])
```

```
corrplot(corMatrix, order = "FPC", method = "color", type = "lower",  
          tl.cex = 0.8, tl.col = rgb(0, 0, 0))
```



### IV Prediction Model Building Three methods will be applied to model the regressions (in the Train dataset) and the best one (with higher accuracy when applied to the Test dataset) will be used for the quiz predictions. The methods are: Random Forests, Decision Tree and Generalized Boosted Model, as described below. A Confusion Matrix is plotted at the end of each analysis to better visualize the accuracy of the models. ##### Method: Random Forest

```
# model fit
set.seed(12345)
controlRF <- trainControl(method="cv", number=3, verboseIter=FALSE)
modFitRandForest <- train(classe ~ ., data=TrainSet, method="rf",
                           trControl=controlRF)
modFitRandForest$finalModel
```

```
##
## Call:
## randomForest(x = x, y = y, mtry = param$mtry)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 27
##
##           OOB estimate of  error rate: 0.23%
## Confusion matrix:
##      A   B   C   D   E class.error
## A 3904    2    0    0    0 0.0005120328
## B    6 2647    4    1    0 0.0041384500
## C    0    5 2391    0    0 0.0020868114
## D    0    0    9 2243    0 0.0039964476
```

```
## E    0    0    0    5 2520 0.0019801980
```

```
# prediction on Test dataset
TestSet$classe=as.factor(TestSet$classe)
predictRandForest <- predict(modFitRandForest, newdata=TestSet)
confMatRandForest <- confusionMatrix(predictRandForest, TestSet$classe)
confMatRandForest
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction    A    B    C    D    E
##           A 1674    1    0    0    0
##           B    0 1138    2    0    0
##           C    0    0 1024    2    0
##           D    0    0    0  962    1
##           E    0    0    0    0 1081
```

```
## Overall Statistics
```

```
##
##           Accuracy : 0.999
##           95% CI : (0.9978, 0.9996)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
##           Kappa : 0.9987
```

```
## McNemar's Test P-Value : NA
```


```
##
## Statistics by Class:
```

```
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.0000  0.9991  0.9981  0.9979  0.9991
## Specificity      0.9998  0.9996  0.9996  0.9998  1.0000
## Pos Pred Value   0.9994  0.9982  0.9981  0.9990  1.0000
## Neg Pred Value   1.0000  0.9998  0.9996  0.9996  0.9998
## Prevalence       0.2845  0.1935  0.1743  0.1638  0.1839
## Detection Rate   0.2845  0.1934  0.1740  0.1635  0.1837
## Detection Prevalence 0.2846  0.1937  0.1743  0.1636  0.1837
## Balanced Accuracy 0.9999  0.9994  0.9988  0.9989  0.9995
```

```
# model fit
set.seed(12345)
modFitDecTree <- rpart(classe ~ ., data=TrainSet, method="class")
fancyRpartPlot(modFitDecTree)
```

Method: Decision Trees

```
## Warning: labs do not fit even at cex 0.15, there may be some overplotting
```



Rattle 2020-jul.-28 19:28:38 victo

*# prediction on Test dataset*

```
predictDecTree <- predict(modFitDecTree, newdata=TestSet, type="class")
confMatDecTree <- confusionMatrix(predictDecTree, TestSet$classe)
confMatDecTree
```

## Confusion Matrix and Statistics

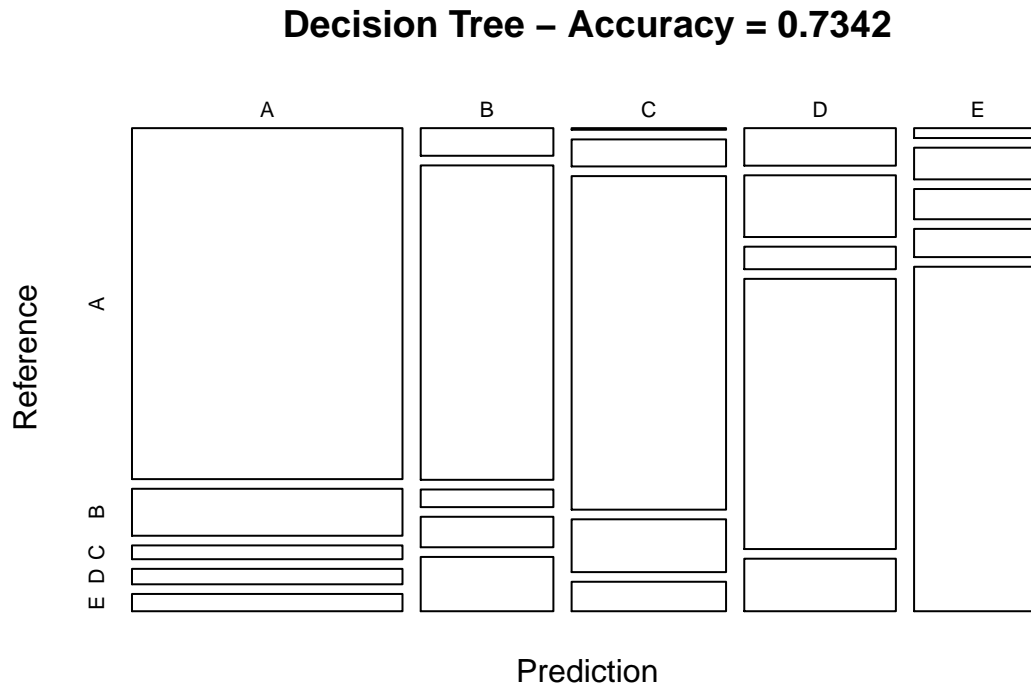
```
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1502  201   59   66   74
##           B   58  660   37   64  114
##           C    4   66  815  129   72
##           D   90  148   54  648  126
##           E   20   64   61   57  696
```

## Overall Statistics

```
##
##           Accuracy : 0.7342
##           95% CI : (0.7228, 0.7455)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6625
##
##           Mcnemar's Test P-Value : < 2.2e-16
##
```

```
## Statistics by Class:
##
##               Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.8973   0.5795   0.7943   0.6722   0.6433
## Specificity      0.9050   0.9425   0.9442   0.9151   0.9579
## Pos Pred Value   0.7897   0.7074   0.7505   0.6079   0.7751
## Neg Pred Value    0.9568   0.9033   0.9560   0.9344   0.9226
## Prevalence       0.2845   0.1935   0.1743   0.1638   0.1839
## Detection Rate    0.2552   0.1121   0.1385   0.1101   0.1183
## Detection Prevalence 0.3232  0.1585   0.1845   0.1811   0.1526
## Balanced Accuracy 0.9011   0.7610   0.8693   0.7936   0.8006
```

```
# plot matrix results
plot(confMatDecTree$table, col = confMatDecTree$byClass,
     main = paste("Decision Tree - Accuracy =",
                  round(confMatDecTree$overall['Accuracy'], 4)))
```



```
#### Method: Generalized Boosted Model
```

```
# model fit
set.seed(12345)
controlGBM <- trainControl(method = "repeatedcv", number = 5, repeats = 1)
modFitGBM <- train(classe ~ ., data=TrainSet, method = "gbm",
                  trControl = controlGBM, verbose = FALSE)
modFitGBM$finalModel
```

```
## A gradient boosted model with multinomial loss function.
```

```
## 150 iterations were performed.
## There were 53 predictors of which 53 had non-zero influence.
```

```
# prediction on Test dataset
```

```
predictGBM <- predict(modFitGBM, newdata=TestSet)
confMatGBM <- confusionMatrix(predictGBM, TestSet$classe)
confMatGBM
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction    A    B    C    D    E
##           A 1668   12    0    1    0
##           B    6 1115   12    1    3
##           C    0   12 1012   21    0
##           D    0    0    2  941    6
##           E    0    0    0    0 1073
```

```
## Overall Statistics
```

```
##
##           Accuracy : 0.9871
##           95% CI : (0.9839, 0.9898)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##           Kappa : 0.9837
```

```
##           McNemar's Test P-Value : NA
```

```
## Statistics by Class:
```

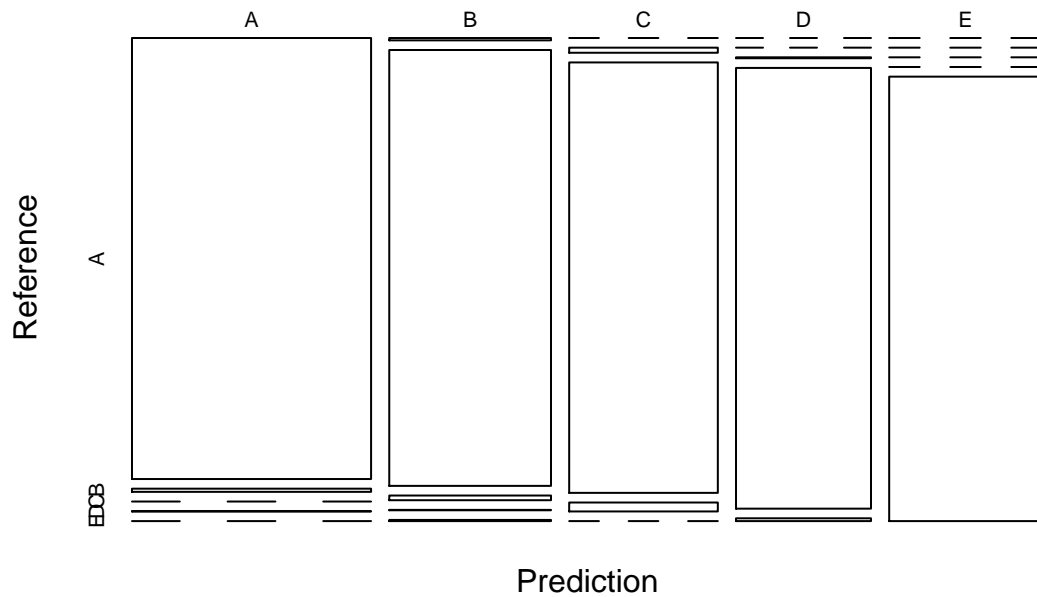
```
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9964  0.9789  0.9864  0.9761  0.9917
## Specificity      0.9969  0.9954  0.9932  0.9984  1.0000
## Pos Pred Value   0.9923  0.9807  0.9684  0.9916  1.0000
## Neg Pred Value   0.9986  0.9949  0.9971  0.9953  0.9981
## Prevalence       0.2845  0.1935  0.1743  0.1638  0.1839
## Detection Rate   0.2834  0.1895  0.1720  0.1599  0.1823
## Detection Prevalence 0.2856  0.1932  0.1776  0.1613  0.1823
## Balanced Accuracy 0.9967  0.9871  0.9898  0.9873  0.9958
```

```
# plot matrix results
```

```
plot(confMatGBM$table, col = confMatGBM$byClass,
     main = paste("GBM - Accuracy =", round(confMatGBM$overall['Accuracy'], 4)))
```



## GBM – Accuracy = 0.9871



### V. Applying the Selected Model to the Test Data The accuracy of the 3 regression modeling methods above are:

1.Random Forest : 0.9963

2.Decision Tree : 0.7368

3.GBM : 0.9839

In that case, the Random Forest model will be applied to predict the 20 quiz results (testing dataset) as shown below.

```
predictTEST <- predict(modFitRandForest, newdata=testing)
predictTEST
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
```

```
## Levels: A B C D E
```